

Bootstrap Learning in Autonomous Intelligent Machines and Systems through Transfer Learning

Ciarán Miceal Johnson

January 2022

1 Topic statement

Yu et al. [10] state that despite the promising ability of Reinforcement Learning (RL) agents to automatically acquire behaviour, the data intensive nature of current approaches precludes the ability to learn multiple tasks, if all tasks are to be learnt individually from scratch. Several fields of Machine Learning (ML) exist which aim to bootstrap learning, two which are of particular interest are: Transfer Learning (TL) and Multitask Learning (MTL). TL focuses on utilising knowledge gained in a source domain to improve the performance of a target learner in a different but similar target domain [14]. MTL on the other hand aims to improve generalisation through an inductive transfer mechanism [3]. Both of these approaches require what is known as positive transfer and suffer from an issue known as Negative Transfer (NT) also referred to as Negative Interference in MTL [12]. NT occurs when utilising knowledge gained in the source domain negatively impacts the ability to learn in the target domain. Similarly, Negative Interference occurs when knowledge previously learnt solving one task negatively affects the ability to learn a new task.

Traditionally RL models require a large dataset of interactions with a chosen environment in order to effectively learn a policy or value function. However due to the time constraints placed on many AIMS, in order for RL to be effective they must learn new tasks from relatively few samples. Utilising prior knowledge may be one path to reducing the experience required to learn. This knowledge can come from the target AIMS itself or learning can be distributed across many AIMS that use a variety of sensors and actuators to interact with the environment. AIMS with different sensors and actuators will have different domains, as how they observe the environment and their ability to interact with it will vary greatly. In order to collaborate effectively TL must be utilised to communicate knowledge learnt in the source domain to the target domain. We may achieve this by exploiting techniques such as domain similarity estimation and distant transfer to help mitigate NT. Solving the problem of NT would improve the ability of distributed AIMS to work together and learn from each other's interactions, reducing the amount of experience a single AIMS needs to collect in order to act effectively in a given environment. This area will likely build upon that of collaborative Multi-Agent Reinforcement Learning (MARL).

The use of TL, to bootstrap learning, could also be beneficial in the field of Embodied Intelligence. Embodied Intelligence addresses the question of how to combine the mind (AI) and the body (robotics) together to work effectively in uncertain environments. This requires rich sensory data and the ability to process large amounts of multi-modal data on edge devices due to the time constraints placed on AIMS. As resource constraints diminish through technological developments, the iterative process of building AIMS with additional improved sensors and actuators will continue. Both of these improvements will alter the domains in which the AIMS operates and therefore it would be wasteful to start the learning process from scratch due to such improvements. Consequently, transferring knowledge gained from one AIMS to another would be of great benefit to Embodied Intelligence. TL may be applicable in this area, though if the new AIMS is too far removed from the previous iterations, then TL could result in NT reducing the ability of the AIMS to learn.

2 Research Goal

My research goal is to bootstrap the learning process of AIMS that interact in uncertain environments by utilising prior knowledge gained from different AIMS. Zhang et al. [12] states that ensuring positive transfer in challenging open environments with continual streams of heterogeneous data is an open area of research in the field of TL. As such I have decided to address this challenge during my PhD. S. J. Pan and Q. Yang [8] categorises the three key challenges in Transfer Learning as what to transfer, how to transfer and when to transfer. The first two challenges have been the subject of extensive research with the third challenge of when to transfer only recently gaining traction in the research community. One of the issues that must be addressed in when to transfer is that of NT. The issue of NT will likely arise during the course of my PhD and is an area that I wish to contribute to in the current research community.

3 Review of the literature

3.1 Negative Transfer

Zhang et al. [12] in their recent survey of NT pull together research from Rosenstein et al. [9] and Pan et al. [8] to establish a set of basic assumptions for TL which when violated can result in NT. These are that

the tasks learnt in both domains must be similar; both domains must have a similar data distribution; and there exists a suitable model which can be applied to both domains. They also group the approaches used to handle NT into four categories: secure transfer, domain similarity estimation, distant transfer and NT mitigation.

Domain estimation concerns itself with computing the transferability between the source and target domain. If domain similarity is low, the transfer of knowledge can be refused or distant TL can be applied. Distant TL, also called transitive TL, bridges the gap between dramatically different domains through multiple intermediate domains. When similarity is high, data from both domains can be concatenated and used to learn in the target domain. When similarity is neither high nor low, NT mitigation can be applied.

Zhu et al. [13] specifies that Modeling Transferability is an issue when performing TL in RL. This issue addresses whether, or to what extent, can knowledge gained from solving one task be applied to solving another. This is akin to domain similarity estimation in TL and therefore solutions used to perform domain similarity estimation in Deep Learning (DL) models may be applicable to Model Transferability in RL. Domain similarity in DL can be estimated through fine-tuning, which is frequently used in deep TL to adapt the source domain DL model to the target domain. The model can be fine-tuned on downstream tasks by fixing lower layer parameters and fine-tuning the higher layer parameters.

MTL solves multiple learning tasks simultaneously by exploiting their commonalities and differences. Similar to TL, it requires positive transfer among tasks to improve the overall learning performance. Previous studies [7, 4] have observed that conflicting gradients among different tasks may induce Negative Interference. Yu et al. [11] proposes PCGrad, a gradient surgery method for MTL that projects task gradients onto the normal plane of all conflicting tasks. PCGrad was able to improve performance and efficiency on various multi-task RL and supervised learning challenges. This approach is useful for avoiding Negative Interference in MTL, though due to the lack of conflicting gradients in single-task TL it will likely be of little benefit when avoiding NT in this scenario.

3.2 Learning from Demonstration

Learning from Demonstration (LfD) looks to use human knowledge to aid in the learning of tasks, this is achieved by imitating experts. Adama et al. [2] discuss how TL in assistive robotics can be effective through the use of expert demonstrators, where the source domain’s features are created by watching an expert, for example a nurse in a care home. The knowledge gained through observing the expert is then transferred to the target domain in which the assistive robot operates. Generative Adversarial Imitation Learning (GAIL) [5] is a model-free imitation learning approach which scales relatively well in high-dimensional environments. This approach draws inspiration from Imitation Reinforcement Learning and Generative Adversarial Networks, requiring few expert interactions to learn expert level policies. The drawback of GAIL is that it has a high sample-complexity during training, meaning it requires an excessive number of interactions with the environment to learn. InfoGAIL [6] is an eXplainable RL (XRL) model which builds upon GAIL by identifying the salient latent factors that differ between demonstrations. XRL and eXplainable AI (XAI) are important areas of research that have been gaining attention recently.

However, relying on demonstrations from human experts to acquire new tasks removes a degree of autonomy from the AIMS. Abbatematteo et al. [1] aims to bootstrap learning using motion planning and therefore avoiding the need for a human demonstrator. Their research uses goal-orientated kinematic planning to generate initial trajectories in order to reduce sample complexity before optimising through policy search. This approach provides a starting point for learning, improving upon learning from scratch with random initialisation yet not outperforming a LfD initialisation. This paper proposes a solution to autonomously bootstrap learning, however it does not address the issue of transferring knowledge from one domain to another in order to bootstrap learning.

4 Research methods and Study Design

The plan for this PhD is to further investigate TL within the fields of RL, MTL, and Embodied Intelligence. Within the field of RL I intended to focus on decentralised cooperative MARL, the goal being to discover techniques used in MARL to communicate effective decentralised policies from one agent to another while ensuring positive transfer. I will also explore the different approaches that MTL uses to avoid Negative Interference in order to determine their applicability in TL.

Upon the completion of my PhD I intend to bootstrap learning in a target AIMS by utilising knowledge gained in a source AIMS. This goal can be broken down into a research phase, where previous research is collated into an academic paper. Following which, I intend to teach my first robot a single task in a deterministic environment, then bootstrap an identical robot so that it can also quickly learn this task. I will then perform MTL on the first robot and again use this to bootstrap learning on the second robot. At this stage I will change the second robot so that it is similar yet different to the first robot, in essence changing the target domain so that it is no longer equivalent to the source domain. It is at this stage I must be wary of NT, and given time I will select a robot whose domain drastically differs from the original robot so that NT is induced. I will then attempt to bootstrap learning in this new, slightly different robot through TL for both the single and multi-task challenge. Finally I intend to deploy the developed methods into an uncertain environment repeating the steps above. Through this research I wish to develop new methodologies to improve TL in uncertain environments and to address issues such as NT and Negative Interference.

References

- [1] Ben Abbatematto, Eric Rosen, Stefanie Tellex, and George Konidaris. Bootstrapping motor skill learning with motion planning, 2021.
- [2] David Adama, Ahmad Lotfi, R. Ranson, and Pedro Trindade. Transfer learning in assistive robotics: From human to robot domain. 01 2019.
- [3] Rich Caruana. Multitask learning. *Machine Learning*, 28(1):41–75, July 1997.
- [4] Zhao Chen, Jiquan Ngiam, Yanping Huang, Thang Luong, Henrik Kretzschmar, Yuning Chai, and Dragomir Anguelov. Just pick a sign: Optimizing deep multitask models with gradient sign dropout, 2020.
- [5] Jonathan Ho and Stefano Ermon. Generative adversarial imitation learning, 2016.
- [6] Yunzhu Li, Jiaming Song, and Stefano Ermon. Infogail: Interpretable imitation learning from visual demonstrations, 2017.
- [7] Bo Liu, Xingchao Liu, Xiaojie Jin, Peter Stone, and Qiang Liu. Conflict-averse gradient descent for multi-task learning, 2021.
- [8] Sinno Jialin Pan and Qiang Yang. A survey on transfer learning. *IEEE Transactions on Knowledge and Data Engineering*, 22(10):1345–1359, 2010.
- [9] Michael Rosenstein, Zvika Marx, Leslie Kaelbling, and Thomas Dietterich. To transfer or not to transfer. 01 2005.
- [10] Tianhe Yu, Saurabh Kumar, Abhishek Gupta, Sergey Levine, Karol Hausman, and Chelsea Finn. Multi-task reinforcement learning without interference. 2019.
- [11] Tianhe Yu, Saurabh Kumar, Abhishek Gupta, Sergey Levine, Karol Hausman, and Chelsea Finn. Gradient surgery for multi-task learning, 2020.
- [12] Wen Zhang, Lingfei Deng, Lei Zhang, and Dongrui Wu. A survey on negative transfer, 2021.
- [13] Zhuangdi Zhu, Kaixiang Lin, and Jiayu Zhou. Transfer learning in deep reinforcement learning: A survey, 2021.
- [14] Fuzhen Zhuang, Zhiyuan Qi, Keyu Duan, Dongbo Xi, Yongchun Zhu, Hengshu Zhu, Hui Xiong, and Qing He. A comprehensive survey on transfer learning. *CoRR*, abs/1911.02685, 2019.